

# Gradient Descent in Gory Detail:

OLS to Mini-batching and Beyond

Hands-on session linear regression + gradient descent

---

Jacob Munson & Will Hammond

Montana State University | TAILS | 10/17/2025 AD

# Learning goals

- Understand ordinary least squares (OLS) for linear regression:
  - Model
  - Loss
  - Closed-form solution
- Gradient descent:
  - Derive gradients of mean-squared error (MSE)
  - Implement gradient descent (GD) updates
- Compare GD flavors: batch GD vs. mini-batch vs. stochastic GD
- Diagnose convergence (learning rate schedules, conditioning, normalization)


# Motivation: Why Linear Regression?

**Scenario:** We observe data pairs  $(x_i, y_i)$  and suspect a linear trend

- Goal: find a line  $\hat{y}_i = \beta_0 + \beta_1 x_i$  that best fits the data
- Minimize Mean Squared Error (MSE):

$$J(\beta_0, \beta_1) = \frac{1}{n} \sum_i (\hat{y}_i - y_i)^2$$

- Captures linear input-output relationships
- Basis for understanding gradient descent
- **In the figure:** the dashed black line is the *true fit* ( $y = 1.5x + 3$ ), the red line is the *OLS fit*, and gray dots are noisy observed data.



**Figure 1:** Simulated data and fitted line.

**Matrix calculus:** simplifies gradients and updates for all parameters simultaneously.

- Let  $X \in \mathbb{R}^{n \times d}$   
(rows are examples, columns are features)
- Response  $y \in \mathbb{R}^n$
- Parameters  $\beta \in \mathbb{R}^d$
- Generalize scalar case  $(\beta_0, \beta_1)$  to vector form:
  - Predictions:  $\hat{y} = X\beta$
  - Loss (MSE):  $J(\beta) = \frac{1}{n} \|X\beta - y\|_2^2$

## Notation summary

symbol	shape	comment
$X$	$n \times d$	design matrix
$y$	$n \times 1$	targets
$\beta$	$d \times 1$	parameters

## OLS: closed-form solution

$$J(\beta) = \frac{1}{n}(X\beta - y)^\top(X\beta - y)$$

$$\nabla_{\beta} J(\beta) = \frac{2}{n}X^\top(X\beta - y)$$

Set gradient to zero:  $X^\top X \beta = X^\top y$

$$\Rightarrow \beta^{\text{OLS}} = (X^\top X)^{-1}X^\top y \quad (\text{if } X^\top X \text{ invertible})$$

## Warm-up: Paper & Pencil View

Given data points  $(x_i, y_i)$ , assume a linear model:

$$\hat{y}_i = \beta_0 + \beta_1 x_i$$

**Our goal:** find  $\beta_0, \beta_1$  minimizing the Mean Squared Error (MSE)

$$J(\beta_0, \beta_1) = \frac{1}{n} \sum_{i=1}^n (\hat{y}_i - y_i)^2$$

**Exercise:**

1. Write down  $J(\beta_0, \beta_1)$  explicitly in terms of  $\beta_0, \beta_1$ .
2. Compute partial derivatives:

$$\frac{\partial J}{\partial \beta_0} = \dots, \quad \frac{\partial J}{\partial \beta_1} = \dots$$

3. Write gradient descent step

# Gradient Descent: Elementwise & Matrix Form

## Elementwise Form

$$\hat{y}_i = \beta_0 + \sum_{j=1}^d x_{ij} \beta_j$$

$$J(\beta) = \frac{1}{n} \sum_i (\hat{y}_i - y_i)^2$$

$$\frac{\partial J}{\partial \beta_j} = \frac{2}{n} \sum_i x_{ij} (\hat{y}_i - y_i)$$

Gradient descent:

$$\beta_j^{(t+1)} = \beta_j^{(t)} - \eta \frac{\partial J}{\partial \beta_j}$$

## Matrix Form

$$\hat{y} = X\beta, \quad J(\beta) = \frac{1}{n} \|X\beta - y\|_2^2$$

Gradient:

$$\nabla_{\beta} J = \frac{2}{n} X^{\top} (X\beta - y)$$

Update rule:

$$\beta^{(t+1)} = \beta^{(t)} - \eta \nabla_{\beta} J$$

**Same math, compactly!**

- Vectorization = efficiency
- Often easier for analysis and coding

# Gradient Descent

**Objective:**  $J(\beta) = \frac{1}{n} \|X\beta - y\|_2^2$

**Gradient:**<sup>1</sup>  $\nabla J(\beta) = \frac{2}{n} X^\top (X\beta - y)$

**Update:**  $\beta^{(t+1)} = \beta^{(t)} - \eta \nabla J(\beta^{(t)})$

## Algorithm (pseudocode)

```
initialize beta randomly or zeros
repeat until convergence:
    grad = (2/n) * X.T @ (X @ beta - y)
    beta = beta - eta * grad
    (should do) monitor J(beta) or validation error
```

---

<sup>1</sup>We often drop the factor of 2  $\rightarrow$  “absorb into  $\eta$ ”



# Gradient Descent Variants

**General update rule:**

$$\beta^{(t+1)} = \beta^{(t)} - \eta \frac{2}{|\mathcal{B}_t|} X_{\mathcal{B}_t}^\top (X_{\mathcal{B}_t} \beta^{(t)} - y_{\mathcal{B}_t})$$

- $\mathcal{B}_t$  is the *batch* of samples used at iteration  $t$ .
- The batch size  $|\mathcal{B}_t|$  determines the variant.

**Special cases:**

- $|\mathcal{B}_t| = n \Rightarrow$  **Batch Gradient Descent**  
Deterministic, stable but slow for large  $n$ .
- $1 < |\mathcal{B}_t| < n \Rightarrow$  **Mini-batch (Stochastic) Gradient Descent**  
Balances efficiency and stability; default in deep learning.
- $|\mathcal{B}_t| = 1 \Rightarrow$  **Stochastic Gradient Descent**  
Highly stochastic updates; rarely used in pure form.

*Note:* In practice, “SGD” often refers to the mini-batch case rather than strictly  $B=1$ .

# Full Algorithm: Mini-batch Gradient Descent

```
# Mini-batch Gradient Descent for Linear Regression
import numpy as np

# X: n x d matrix, y: n-vector
n, d = X.shape
beta = np.zeros(d)
eta = 0.1 # learning rate
B = 32 # batch size
n_epochs = 100

for epoch in range(n_epochs):
    perm = np.random.permutation(n)
    X, y = X[perm], y[perm]
    for i in range(0, n, B):
        idx = slice(i, i+B)
        Xb, yb = X[idx], y[idx]
        grad = (2/B) * Xb.T @ (Xb @ beta - yb)
        beta -= eta * grad

print("Learned coefficients:", beta)
```

## Ridge (L2) regularization

Objective:  $J_\lambda(\beta) = \frac{1}{n}\|X\beta - y\|^2 + \lambda\|\beta\|^2$ .

Closed form:  $\beta = (X^\top X + n\lambda I)^{-1}X^\top y$ .

Gradient:  $\nabla J_\lambda = \frac{2}{n}X^\top(X\beta - y) + 2\lambda\beta$ .

Effect: shrinks coefficients, improves conditioning of  $X^\top X$ ; GD tolerates larger  $\eta$ .

# Compare OLS vs GD variants

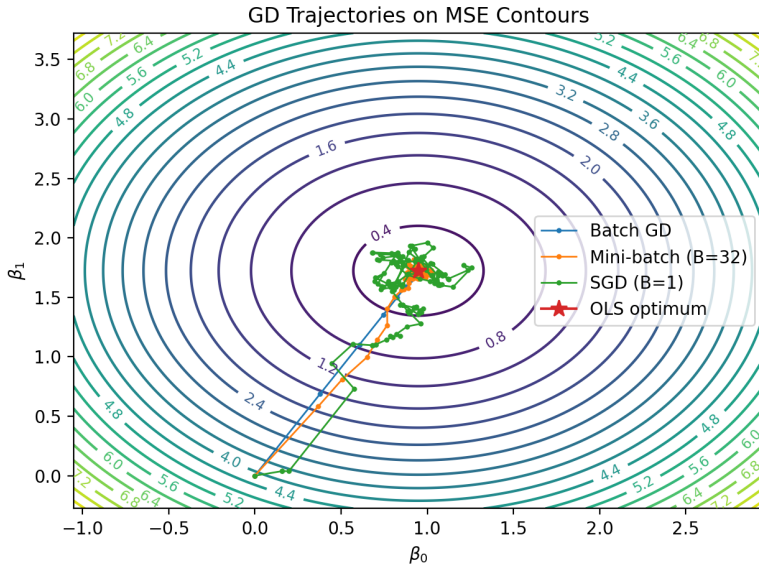
## Metrics

- Parameter error  $\|\beta_{\text{est}} - \beta_{\text{OLS}}\|$ .
- Train/validation MSE.
- Time-to- $\varepsilon$  and epochs to converge.
- Sensitivity to feature scaling.

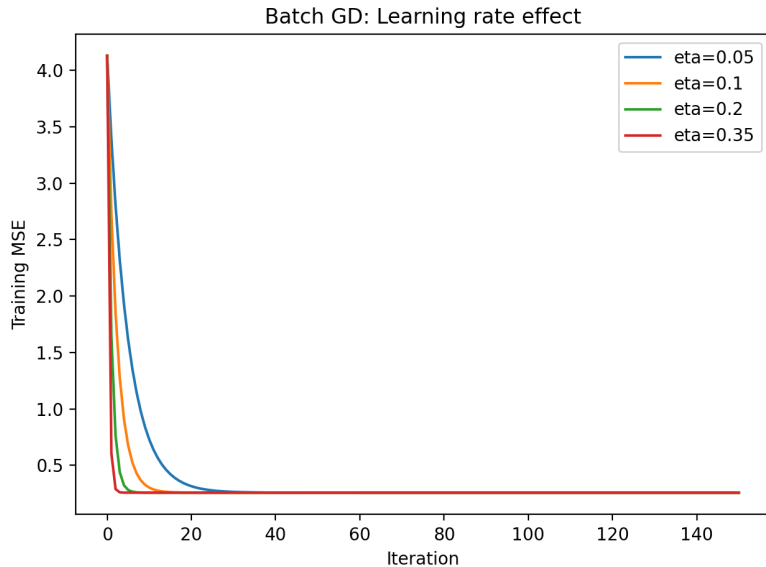
## Why pick one over another?

- OLS is one-shot but unstable for ill-conditioned/huge  $X$ .
- GD scales to large  $n, d$  and streaming data.
- Mini-batch leverages hardware; SGD can escape shallow minima in nonconvex problems.


# GD trajectories on MSE contours




# Learning rate effect (batch GD)



# Batch size vs. convergence noise

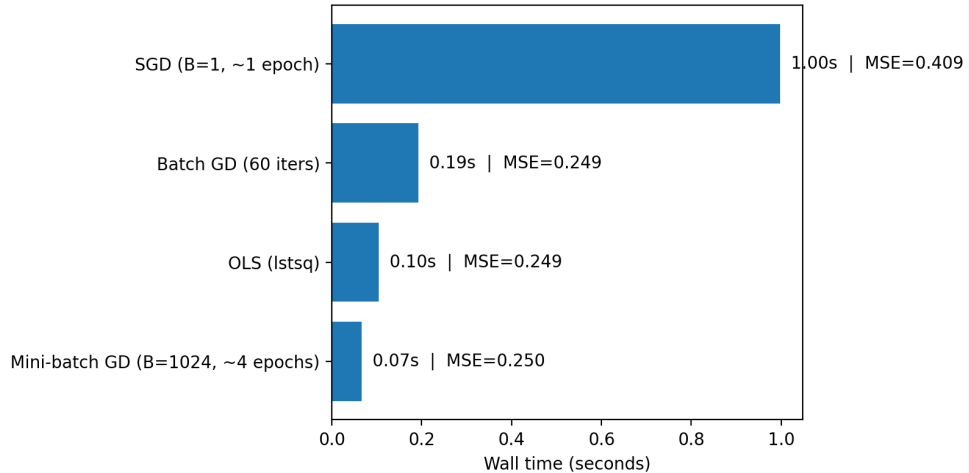


# Parameter convergence toward OLS





Runtime vs. Fit Quality  
Bars: wall-clock time; Labels: seconds and training MSE



# Complexity and Runtime Takeaways

## Computational Complexity

- **Closed-form OLS:**  $\mathcal{O}(nd^2 + d^3)$  time
- **Batch Gradient Descent:**  $\mathcal{O}(nd)$  per iteration
- **Mini-batch / SGD:**  $\mathcal{O}(Bd)$  per step

## Runtime Insights

- **OLS:** One-shot and very fast for small  $d$  and moderate  $n$  (exact up to FP precision)
- **GD/Mini-batch:** Preferred for large  $n$  or streaming data
- **Practical tuning:** Tune learning rate  $\eta$ ; monitor wall time and validation error.

# Common pitfalls

- Forgetting intercept column or feature scaling.
- Learning rate too big (divergence) or too small (stalling).
- Not shuffling mini-batches; data order bias.
- Evaluating only training MSE; always hold out validation.

**Thanks! Questions + discussion.**